Development pilot: clean and retry retrieval contexts

Every cell shows clean above retry. Distinct events are counts; annotated support is percent (rounded).

Lower after retry | Unchanged | Higher after retry | Unavailable support

	LoCoMo A (6)						LoCoMo B (6)						LongMemEval (12)											
	Q01	Q02	Q03	Q04	Q05	Q06	Q07	Q08	Q09	Q10	Q11	Q12	Q13	Q14	Q15	Q16	Q17	Q18	Q19	Q20	Q21	Q22	Q23	Q24
Mem0 raw Distinct events 20/24 lower	10 6	10 8	10 6	10 6	10 5	10 8	10 6	10 10	10 8	10 8	10 6	10 6	10 10	10 8	8 6	10 6	8 8	8 6	10 6	10 10	9 4	10 9	10 4	10 8
Mem0 raw Annotated support % 1/23 lower; 1 NA	100 100	100 100	100 100	100 100	100 100	67 67	100 0	100 100	100 100	100 100	0 0	100 100	100 100	0 0	100 100	100 100	100 100	25 25	NA NA	100 100	100 100	100 100	100 100	100 100
LlamaIndex Distinct events 16/24 lower	10 4	10 7	10 6	10 5	10 6	10 6	10 7	10 10	10 8	10 7	10 10	10 6	10 10	10 10	10 7	10 5	10 10	10 10	10 6	10 10	10 4	10 10	10 4	10 6
LlamaIndex Annotated support % 4/23 lower; 1 NA	100 0	100 100	100 100	0 0	50 0	0 0	100 0	0 0	100 100	100 100	0 0	100 100	100 100	100 100	100 100	100 100	100 100	50 50	NA NA	100 100	100 100	100 100	100 50	100 100

Questions share two LoCoMo conversations or LME evidence components; LME background sessions can overlap.

Common limits: 60 candidates; 10 records; 12,000 serialized evidence bytes. No reader-accuracy or harm claim.

Unavailable support stays NA. Gold annotations may be incomplete. Descriptive paired pilot; no p-values or CI.